

# Image Fragment Reconstruction via Adjacency Prediction

May 3, 2026

## Contents

|           |                                                     |           |
|-----------|-----------------------------------------------------|-----------|
| <b>1</b>  | <b>Problem Statement</b>                            | <b>2</b>  |
| <b>2</b>  | <b>Related Work</b>                                 | <b>2</b>  |
| <b>3</b>  | <b>Core Idea</b>                                    | <b>3</b>  |
| <b>4</b>  | <b>Architecture: Siamese Network</b>                | <b>3</b>  |
| <b>5</b>  | <b>Training</b>                                     | <b>5</b>  |
| 5.1       | Constructing Training Labels . . . . .              | 5         |
| 5.2       | Loss Function . . . . .                             | 5         |
| 5.3       | Computational Cost . . . . .                        | 6         |
| <b>6</b>  | <b>From Pretext Task to Clustering</b>              | <b>6</b>  |
| <b>7</b>  | <b>Evaluation Metric for the Clustering Task</b>    | <b>7</b>  |
| <b>8</b>  | <b>Experimental Setup</b>                           | <b>8</b>  |
| <b>9</b>  | <b>Results</b>                                      | <b>9</b>  |
| 9.1       | Effect of the loss tilt parameter $\beta$ . . . . . | 9         |
| 9.2       | Headline numbers . . . . .                          | 10        |
| <b>10</b> | <b>Failure Modes</b>                                | <b>10</b> |
| <b>11</b> | <b>Decision Points Summary</b>                      | <b>10</b> |
| <b>12</b> | <b>Limitations</b>                                  | <b>11</b> |
| <b>13</b> | <b>Further Suggestions</b>                          | <b>12</b> |

# 1 Problem Statement

Let  $\mathcal{I} = \{I_1, \dots, I_N\}$  be a set of  $N$  source images, each of spatial resolution  $H \times W \times 3$ . Each image is divided into a regular  $G \times G$  grid of non-overlapping fragments of size  $\frac{H}{G} \times \frac{W}{G} \times 3$ , yielding  $G^2$  fragments per image and  $M = N \cdot G^2$  fragments in total. In the experimental setting used here  $N = 10$ ,  $H = W = 64$ ,  $G = 4$ , giving fragments of size  $16 \times 16 \times 3$  and  $M = 160$  fragments per batch.

Let  $f_k \in \mathbb{R}^{16 \times 16 \times 3}$  denote fragment  $k$ , and let  $s(k) \in \{1, \dots, N\}$  be its source image. The fragments from all  $N$  images are pooled into a single unordered collection. The task is:

*Given the pooled collection of  $M$  fragments with no positional or source information, assign each fragment to one of  $N$  groups such that all fragments assigned to the same group originate from the same source image.*

Formally, the goal is to recover a partition  $\hat{\mathcal{C}} = \{\hat{C}_1, \dots, \hat{C}_N\}$  of  $\{1, \dots, M\}$  such that for each predicted cluster  $\hat{C}_n$  there exists a source image  $I_s$  with  $s(k) = s$  for all  $k \in \hat{C}_n$ .

**Constraints.** Two constraints make this non-trivial:

1. **No positional information.** The original grid coordinates of each fragment are not provided at inference time. The model may not use absolute position as a feature.
2. **No labels during training.** The source identity  $s(k)$  is never provided to the model as a supervisory signal. The model must learn from structural properties of the fragments themselves.

**Known structure.** Although labels are withheld, the following quantities are known and may be exploited:

- The total number of source images  $N = 10$  (so  $k = N$  clusters are sought).
- Each source image contributes exactly  $G^2 = 16$  fragments, so all clusters have equal size.
- Spatial adjacency within a source image is fully determined by the grid structure and can be derived without labels: fragments at grid positions  $(r, c)$  and  $(r', c')$  are adjacent if and only if  $|r - r'| + |c - c'| = 1$ .

The third point is the key observation exploited by the proposed approach (Section 2).

## 2 Related Work

The problem sits at the intersection of self-supervised representation learning and clustering. We review three families of prior work that informed the approach.

**Pretext-task self-supervision.** The dominant paradigm in self-supervised vision until the rise of contrastive methods was to define a hand-crafted pretext task whose labels are derivable from the data itself. Noroozi and Favaro (2016) solve a jigsaw permutation: given a  $3 \times 3$  grid of fragments shuffled by one of 100 fixed permutations, predict which permutation was used. The encoder learns features useful for downstream classification. Our adjacency prediction is a much simpler, binary variant of this idea — rather than predicting a discrete permutation index over an exponentially large space, we predict a binary label per pair. This trades off representational richness for sample efficiency.

**Contrastive learning.** SimCLR (Chen et al. 2020) and MoCo (He et al. 2020) learn representations by pulling together augmented views of the same image and pushing apart views of different images. In the present setting these methods are awkward: there is no natural positive pair construction, since the natural “view” of a fragment is the fragment itself. One could treat all fragments from one source image as positives of each other, but this requires the source label that we are explicitly forbidden to use during training. Adjacency labels, in contrast, are derivable from the grid structure alone.

**Clustering-based self-supervision.** DeepCluster (Caron, Bojanowski, et al. 2018) alternates between clustering current embeddings and using cluster assignments as pseudo-labels. SwAV (Caron, Misra, et al. 2020) extends this with online clustering and equipartitioned cluster assignments via the Sinkhorn algorithm. Both are conceptually well-suited to this task, but they introduce significant complexity (alternating training, prototype maintenance) for what is at heart a small-scale clustering problem with  $N = 10$  images. We treat this family as a strong candidate for follow-up work but defer it.

### 3 Core Idea

The central observation is that two fragments cut from the same image and placed next to each other in the grid share a common boundary — colours bleed over, textures continue, and objects span tiles. This visual continuity is a strong and reliable signal that does not require any human annotation to exploit: since the experimenter creates the grid, the ground-truth adjacency of every fragment pair is known at all times.

The proposed approach uses adjacency prediction as a **pretext task** (Noroozi and Favaro 2016). The model is trained to solve a surrogate problem — predicting whether two fragments are spatially adjacent — and in doing so is forced to learn visual representations that capture source identity as a side effect. At inference time the pretext objective is discarded and the learned representations are used for clustering.

This is self-supervised throughout: no labels beyond those derivable from the grid structure are ever used.

### 4 Architecture: Siamese Network

The model is a Siamese network: a shared encoder is applied to each of two input fragments, and a comparison head maps the two embeddings to a confidence score  $p_{ij} \in [0, 1]$  for the binary question of interest (Figure 1). When the multi-task loss is enabled, a second comparison head with the same structure is added for the same-image task.

The components are:




Figure 1: Schematic of the model. The CNN encoder is applied with shared weights to each of two input fragments; its 256-dimensional embeddings are combined into a single 1024-dimensional vector via concatenation of difference, element-wise product, and the two embeddings themselves; this vector is fed to one or two comparison heads (adjacency, same-image) that produce a confidence score  $p_{ij} \in [0, 1]$  for the corresponding binary question.

**Shared CNN encoder.** A single convolutional neural network maps each  $16 \times 16 \times 3$  fragment to a 256-dimensional embedding vector. The architecture is three convolutional blocks (Conv–BatchNorm–ReLU), with max-pooling after blocks two and three, reducing the spatial dimensions from  $16 \times 16$  to  $4 \times 4$ . The flattened  $4 \times 4 \times 128$  feature map is projected to a 256-dimensional embedding via a linear layer followed by ReLU and dropout ( $p = 0.3$ ).

The encoder is **shared** — the same weights are used for both fragments in a pair. This is the defining property of a Siamese network (Bromley et al. 1994): weight sharing ensures that the comparison is symmetric and that the encoder learns a universal fragment representation rather than one tuned to a specific input slot.

**Comparison head (MLP).** Let  $\mathbf{a} = \text{CNN}(f_i)$  and  $\mathbf{b} = \text{CNN}(f_j)$  denote the two embeddings produced by the shared encoder for an input pair. The comparison head takes as input the 1024-dimensional vector

$$\mathbf{z}_{ij} = [\mathbf{a} - \mathbf{b} \parallel \mathbf{a} \odot \mathbf{b} \parallel \mathbf{a} \parallel \mathbf{b}] \in \mathbb{R}^{1024},$$

formed by concatenating ( $\parallel$ ) four sub-vectors of size 256: the element-wise difference  $\mathbf{a} - \mathbf{b}$ , capturing direction and magnitude of dissimilarity; the element-wise (Hadamard) product  $\mathbf{a} \odot \mathbf{b}$ , capturing feature co-activation (large only when both fragments score high on the same feature); and the two raw embeddings, preserving each fragment’s individual information. This four-term combination is borrowed from the InferSent sentence-pair classification model of Conneau et al. (2017), where it was shown to outperform plain concatenation on natural-language inference. We adopt it here without ablation in our setting; whether the explicit difference and product terms help on image fragments, or whether a simpler  $[\mathbf{a}, \mathbf{b}]$  concatenation with a larger hidden layer would do as well, is left as future work. The MLP maps  $\mathbf{z}_{ij}$  through that hidden layer (ReLU) to a single scalar in  $[0, 1]$ :

$$p_{ij} = \sigma(W_2 \text{ReLU}(W_1 \mathbf{z}_{ij} + b_1) + b_2).$$

## 5 Training

### 5.1 Constructing Training Labels

From a sample of 10 images cut into a  $4 \times 4$  grid, each fragment has at most 4 adjacent neighbours (fewer at corners and edges). For a sample of 160 fragments the total number of pairs is:

$$\binom{160}{2} = 12,720$$

Each pair receives a binary label:  $y_{ij} = 1$  if the two fragments are adjacent within the same image,  $y_{ij} = 0$  otherwise. This label is derived entirely from the known grid positions and requires no human annotation (Figure 2).




Figure 2: Adjacency labelling within a single source image. Each of the 16 fragments is connected by a red edge to each of its grid neighbours. The 24 highlighted edges are the positive pairs ( $y_{ij} = 1$ ); all other pairs within the image (96 of them) and across images are negative ( $y_{ij} = 0$ ).

### 5.2 Loss Function

Binary cross-entropy is the natural loss for a binary prediction task:

$$\mathcal{L}_{\text{BCE}} = -\frac{1}{N} \sum_{(i,j)} \left[ y_{ij} \log p_{ij} + (1 - y_{ij}) \log(1 - p_{ij}) \right]$$

Positive (adjacent) pairs are far rarer than negative pairs in this configuration. With 10 source images on a  $4 \times 4$  grid, each image contributes  $4 \cdot 3 + 3 \cdot 4 = 24$  adjacent pairs, giving  $n_+ = 240$  positives in total. The remaining pairs are negatives:  $n_- = \binom{160}{2} - 240 = 12,480$ . The ratio  $n_-/n_+ = 52$  is therefore a constant determined by the grid geometry, not a per-batch quantity. Vanilla BCE on this data is dominated by the negative class: the gradient is overwhelmingly driven by the loss on non-adjacent pairs, and the model has little incentive to identify the few adjacent ones. The standard remedy is to upweight the rare class. Since only the ratio of the two class weights matters, we parameterise the loss with a single *tilt* parameter  $\beta \geq 1$ :

$$\mathcal{L}_{\text{WBCE}}(\beta) = -\frac{1}{N} \sum_{(i,j)} \left[ \beta \cdot \frac{n_-}{n_+} y_{ij} \log p_{ij} + (1 - y_{ij}) \log(1 - p_{ij}) \right].$$

The textbook value is  $\beta = 1$ , the class-balanced loss: it equalises the total gradient contribution of the two classes, treating a single positive as worth 52 negatives. Setting  $\beta > 1$  over-weights positives beyond this point — a positive mistake is now *more* costly than the natural class ratio implies. For our downstream task this is desirable, and we expect the optimum to lie at a moderate-to-aggressive tilt.

**Why we lean toward a moderate-to-aggressive tilt.** The downstream task is graph clustering, not per-edge classification. False negatives directly weaken intra-cluster connectivity, while false positives are only harmful when they connect fragments from *different* source images — a case the model rarely produces because cross-image fragments are visually dissimilar. Same-image false positives, the more common kind, actually reinforce the connectivity spectral clustering relies on. The error costs are therefore asymmetric in favour of recall, and we expect the optimal  $\beta$  to lie above the class-balanced point.

### 5.3 Computational Cost

The CNN encoder runs once per fragment (160 forward passes per sample). All pairwise embeddings are then combined using tensor operations, and the lightweight MLP processes all 12,720 pairs. Since the expensive step (the CNN) runs only 160 times, the computational overhead is manageable.

## 6 From Pretext Task to Clustering

After training, the model produces a value  $p_{ij} \in [0, 1]$  for every pair of fragments. We do not interpret this as a calibrated probability — the model is trained on a binary cross-entropy objective but its outputs are not guaranteed to be calibrated to true frequencies. We use  $p_{ij}$  instead as a confidence score: the higher the value, the more confident the model is that fragments  $i$  and  $j$  are adjacent. These scores define a **fully connected weighted graph** over the 160 fragments, where edge weight  $w_{ij} = p_{ij}$  reflects the model’s confidence that the two fragments are adjacent — and by extension, that they originate from the same image (Figure 3).

**Balanced spectral clustering.** **Spectral clustering** is the natural choice for a weighted graph. Given the  $160 \times 160$  similarity matrix  $W$  with  $w_{ij} = p_{ij}$ , define the diagonal degree matrix  $D$  with  $D_{ii} = \sum_j w_{ij}$  and the symmetric normalised graph Laplacian

$$L_{\text{sym}} = I - D^{-1/2} W D^{-1/2}.$$

The spectral embedding stacks the  $k$  leading eigenvectors of  $L_{\text{sym}}$  into a  $160 \times k$  matrix; in the eigenspace, fragments belonging to densely-connected sub-graphs lie close together. The standard pipeline then runs  $k$ -means on this embedding to obtain hard cluster assignments.

Plain  $k$ -means ignores the fact that each true cluster contains exactly  $G^2 = 16$  fragments and can produce groups of very unequal size. We enforce balance by replacing the  $k$ -means step with a balanced assignment: build a  $160 \times (k \cdot G^2)$  cost matrix in which each cluster appears  $G^2$  times, and solve the resulting linear assignment problem with




Figure 3: Schematic of the predicted similarity graph. Each node is a fragment; edge weight  $p_{ij}$  is the model’s confidence that the two fragments are adjacent. Edges within a true source image (here, two illustrative source images shown in blue and orange) tend to receive high weight; cross-image edges receive low weight. Spectral clustering recovers the source-image partition by finding the cut that separates the densely-connected sub-graphs.

the Hungarian algorithm<sup>1</sup>. Cluster centroids are then re-estimated and the assignment iterated until convergence (typically a handful of iterations). The result is a partition of the 160 fragments into 10 clusters of size 16, by construction.

## 7 Evaluation Metric for the Clustering Task

The pretext task (adjacency prediction) and the downstream task (clustering of fragments by source image) are evaluated with different metrics. Adjacency prediction is evaluated with the threshold-free metrics AUROC and AUPRC (described below in the multi-batch averaging paragraph); the downstream clustering is evaluated with the **Adjusted Rand Index (ARI)**. We focus on the clustering metric here since it is the headline number for this work.

**Definition.** Let  $\mathcal{C} = \{C_1, \dots, C_N\}$  be the true partition of the fragments into  $N = 10$  source-image groups, and  $\hat{\mathcal{C}} = \{\hat{C}_1, \dots, \hat{C}_N\}$  the partition predicted by the model. For each pair of fragments, look at whether the two are in the same true group and whether they are in the same predicted cluster:

- $a$  = number of pairs in the same true group *and* the same predicted cluster
- $b$  = number of pairs in different true groups *and* different predicted clusters
- $c$  = number of pairs in the same true group but different predicted clusters
- $d$  = number of pairs in different true groups but the same predicted cluster

The (unadjusted) Rand Index is the agreement rate  $RI = (a + b)/(a + b + c + d)$ . The Adjusted Rand Index normalises this against the expected RI under random labellings of

<sup>1</sup>Implemented using `sklearn.manifold.SpectralEmbedding` for the eigendecomposition and `scipy.optimize.linear_sum_assignment` for the assignment.

fixed cluster sizes:

$$\text{ARI} = \frac{\text{RI} - \mathbb{E}[\text{RI}]}{\max(\text{RI}) - \mathbb{E}[\text{RI}]}.$$

A random clustering scores  $\text{ARI} = 0$ , a perfect one scores  $\text{ARI} = 1$ , and a clustering worse than random can score below zero.

**Why ARI rather than purity.** ARI is preferred over simpler metrics like purity because purity only looks at the dominant true class in each predicted cluster and ignores how the remaining errors are distributed.

**Example.** Suppose two predicted clusters each have 60% fragments from image A. The first additionally has 40% from image B; the second has 20% from B and 20% from C. Both clusters score the same purity (0.60), even though the second is visibly more disordered. ARI distinguishes the two cases because it counts pair-level agreements rather than looking only at the plurality label.

**NMI (Normalised Mutual Information)** is reported alongside ARI as a complementary metric. It measures how much information the predicted clusters share with the true groupings, normalised to  $[0, 1]$ .

**Multi-batch averaging.** Each evaluation step samples a fresh batch of 10 source images from the validation generator. Because individual batches vary substantially in difficulty (some draws of 10 ImageNet classes are visually well-separated, others contain near-duplicates), a single-batch ARI estimate has high variance. We average each metric over  $n_{\text{eval}} = 20$  fresh validation batches per evaluation step. This stabilises the training curves and, crucially, the early-stopping decision and the checkpoint-selection criterion, which would otherwise be biased toward lucky batches.

## 8 Experimental Setup

**Data.** ImageNet-64 (Chrabaszcz, Loshchilov, and Hutter 2017): a downsampled version of the ImageNet-1k classification dataset in which every image has been resized to  $64 \times 64$  pixels. The dataset is split into a training partition and a test partition; we draw the model’s training images from the former and the validation images from the latter. Each training step samples 10 random images on the fly and produces a fresh batch of 160 fragments; the data generator is therefore effectively infinite. The augmentation pipeline (rotation up to  $\pm 20^\circ$ , horizontal/vertical shift up to 10%, channel shift 0.2) is applied to the full image *before* fragmentation and is provided by the dataset’s data generator (`data.py`); it is not modified.

**Hardware.** All training and evaluation reported in this document was performed on a MacBook Air (Apple M-series CPU). PyTorch is configured to fall back to CPU; no GPU is required. A full training run completes in approximately 20–25 minutes on this hardware.

## Reproducibility.

- Each run writes its full configuration, training log, and best checkpoint into `runs/{run_name}/`, and appends a one-line summary to `results.jsonl`. The configuration is therefore self-contained and re-runnable.
- Random seeds are fixed for the clustering routines<sup>2</sup>. The PyTorch random seed is *not* currently fixed, so model initialisation and the order of training batches vary between runs; this is by design — repeated runs with different seeds give a sense of the variance in final ARI.
- Hyperparameters and run metadata are versioned in `config.yaml`.

**Hyperparameters.** The complete list of hyperparameters is maintained in `config.yaml` and may still change in the course of further ablations. The current values are documented there; this section will be replaced with a final hyperparameter table once the study is complete.

## 9 Results

### 9.1 Effect of the loss tilt parameter $\beta$

We evaluated three settings of the WBCE tilt parameter —  $\beta \in \{0.019, 0.3, 1.0\}$ , corresponding respectively to plain unweighted BCE ( $\beta \cdot 52 = 1$ ), a mild upweighting of positives, and the textbook class-balanced weighting ( $w_+ = 52$ ). All three runs use the same architecture, balanced spectral clustering, and validation averaging over  $n_{\text{eval}} = 20$  fresh batches.

| Run          | $\beta$ | Best ARI | AUROC / AUPRC |
|--------------|---------|----------|---------------|
| plain_bce    | 0.019   | 0.535    | 0.969 / 0.432 |
| wbce_beta_03 | 0.3     | 0.523    | 0.973 / 0.435 |
| wbce_beta_1  | 1.0     | 0.462    | 0.969 / 0.405 |

**Discussion.** The choice of  $\beta$  has an essentially negligible effect on AUROC: across two orders of magnitude in  $\beta$ , ranking quality on the adjacency task varies by less than a percentage point (0.969–0.973). What changes with  $\beta$  is the model’s *calibration* — its tendency to push probabilities high or low overall — not the order in which it ranks pairs. Because spectral clustering operates on the raw probabilities as edge weights, both ranking and calibration matter for ARI, which is why ARI varies a little more than AUROC (0.46–0.54).

The variation we observe in ARI is, however, on the same order as the run-to-run noise we estimated by repeating  $\beta = 0.3$  and  $\beta = 1$  with different random seeds: between repeated runs of the same configuration the ARI changed by up to  $\sim 0.05$ , comparable to the differences between configurations. The cleanest reading of the data is that *within the range tested,  $\beta$  is a second-order knob*. Plain BCE performs at least as well as any weighted variant we tried.

---

<sup>2</sup>`random_state = 42` in `sklearn.cluster.SpectralClustering`, `sklearn.cluster.KMeans`, and `sklearn.manifold.SpectralEmbedding`.

This contradicts the prior expectation expressed earlier in the document that an aggressive tilt  $\beta \in [1.5, 3]$  would be optimal. The argument about asymmetric error costs (false negatives breaking intra-cluster connectivity more than false positives) is not wrong in principle, but the model already learns a near-optimal ranking under plain BCE, so the loss weighting has little remaining effect on the downstream task. We adopt plain BCE for simplicity in subsequent experiments.

## 9.2 Headline numbers

|                                        | Best ARI | AUROC |
|----------------------------------------|----------|-------|
| Run 1 (single-batch eval, unbalanced)  | 0.787    | –     |
| Plain BCE (multi-batch eval, balanced) | 0.535    | 0.969 |

The apparent regression from 0.787 to 0.535 is almost entirely the multi-batch averaging unmasking the variance of single-batch evaluation: the original number was a single lucky draw, while the new number is the mean over twenty fresh validation batches.

**Cross-run comparison.** *(Placeholder for the markdown table produced by `compare_runs.py`, plus the overlay plot `runs_comparison.png`.)*

## 10 Failure Modes

Figure 4 shows four misclassified fragments from the best checkpoint. The model produces only an unordered partition into  $k = 10$  clusters with arbitrary integer IDs; to relate them to source images for this visualisation we align cluster IDs to true source IDs by Hungarian matching on the validation batch (`scipy.optimize.linear_sum_assignment` on the  $10 \times 10$  confusion matrix). The rightmost column then shows the source image whose fragments form the plurality of the cluster the misclassified fragment was assigned to. The matching is per-batch and uses the true labels; the model itself has no notion of source identity.

A reliable taxonomy of failure modes would require aggregating misclassified fragments across many batches, which we leave as follow-up work.

## 11 Decision Points Summary

| Decision              | Choice                                | Alternatives               |
|-----------------------|---------------------------------------|----------------------------|
| Training objective    | Adjacency prediction (pretext task)   | Contrastive loss           |
| Architecture          | Siamese CNN + MLP head                | Single encoder             |
| Embedding combination | Diff + product + both embeddings      | Concatenation only         |
| Loss function         | Binary cross-entropy                  | Focal loss (for imbalance) |
| Class imbalance       | Weighted loss or negative subsampling | None                       |
| Clustering algorithm  | Balanced spectral clustering          | Plain spectral, $k$ -means |
| Number of clusters    | $k = 10$ (known)                      | Discovered automatically   |
| Cluster size          | Enforced ( $G^2 = 16$ ) via Hungarian | Free                       |
| Eval batches per step | 20 (averaged)                         | 1 (single batch)           |
| Evaluation metric     | ARI, NMI                              | Purity                     |




Figure 4: Selected misclassified fragments from the best checkpoint (`multitask_10`, best ARI = 0.570,  $\beta = 0.01923$ ,  $\lambda_{\text{adj}} = 1.0$ ,  $\lambda_{\text{same}} = 1.0$ ). On the validation batch shown, 51 of 160 fragments were assigned to the wrong cluster after Hungarian matching of predicted to true labels. Each row shows: the misclassified fragment (left, red), the true source image (middle, green), and the source image the model assigned it to (right, orange).

## 12 Limitations

- Adjacency is a local signal — it says nothing directly about fragments that are far apart within the same image. The model must generalise from local adjacency to global source identity.
- The graph has 12,720 edges. For larger grids or larger samples this grows quadrat-

ically and may become a bottleneck.

- The pretext metric and the task metric are misaligned. We train binary adjacency (only  $\sim 1.9\%$  positives) but evaluate clustering. The model can achieve a useful similarity matrix ( $\text{ARI} \approx 0.78$ ) while only mediocre on adjacency F1 ( $\approx 0.26$ ), because adjacency is a sufficient but not necessary condition for same-source identity.

## 13 Further Suggestions

**Differentiable surrogate for ARI.** The training loss is on adjacency (or same-image), not on the downstream clustering objective. A more direct alternative is to optimise a differentiable approximation of ARI: the eigendecomposition and cluster-assignment steps of spectral clustering can be relaxed via Sinkhorn iterations (as in SwAV) or via a MinCut-style differentiable graph pooling. This would couple the representation learning to the clustering objective explicitly. We did not pursue it because the multi-task same-image head already captures most of the pairwise structure that ARI ultimately depends on, at a fraction of the implementation cost.

**Self-supervised pretrained encoder.** DINOv2 (He et al. 2020) or another self-supervised image encoder could replace our trained-from-scratch CNN. Since these encoders use no labels, this stays within the constraints of the task. Whether the gain (richer features) outweighs the loss (features were trained at  $224 \times 224$ , not  $16 \times 16$ ) is an open question.

**Topology-aware clustering.** The true partition is more constrained than “equal sizes”: each cluster’s induced subgraph is isomorphic to a  $G \times G$  grid (16 nodes, 24 edges, a planar graph with maximum degree 4). A clustering algorithm that exploited this structure, e.g. a small GNN trained to find equipartitions whose induced subgraphs are grid-like, could disambiguate borderline assignments that pure spectral clustering misses. Implementation cost is significant; the marginal benefit beyond what the model’s already-strong soft signal recovers is unclear.

**Color-histogram baseline.** A non-learned baseline using per-fragment color histograms with chi-square or EMD distance, followed by spectral clustering, would anchor the learned model’s ARI against a meaningful reference. We expect this to be competitive on uniform-region fragments and weak on textured ones; a controlled comparison would quantify how much of the 0.57 ARI comes from learned representation vs. trivial color matching.

## References

- Bromley, Jane et al. (1994). “Signature Verification Using a Siamese Time Delay Neural Network”. In: *Advances in Neural Information Processing Systems (NeurIPS)*. Vol. 6, pp. 737–744.
- Caron, Mathilde, Piotr Bojanowski, et al. (2018). “Deep Clustering for Unsupervised Learning of Visual Features”. In: *European Conference on Computer Vision (ECCV)*, pp. 132–149.

- Caron, Mathilde, Ishan Misra, et al. (2020). “Unsupervised Learning of Visual Features by Contrasting Cluster Assignments”. In: *Advances in Neural Information Processing Systems (NeurIPS)*. Vol. 33, pp. 9912–9924.
- Chen, Ting et al. (2020). “A Simple Framework for Contrastive Learning of Visual Representations”. In: *International Conference on Machine Learning (ICML)*, pp. 1597–1607.
- Chrabaszcz, Patryk, Ilya Loshchilov, and Frank Hutter (2017). “A Downsampled Variant of ImageNet as an Alternative to the CIFAR Datasets”. In: *arXiv preprint arXiv:1707.08819*.
- Conneau, Alexis et al. (2017). “Supervised Learning of Universal Sentence Representations from Natural Language Inference Data”. In: *Empirical Methods in Natural Language Processing (EMNLP)*, pp. 670–680.
- He, Kaiming et al. (2020). “Momentum Contrast for Unsupervised Visual Representation Learning”. In: *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 9729–9738.
- Noroozi, Mehdi and Paolo Favaro (2016). “Unsupervised Learning of Visual Representations by Solving Jigsaw Puzzles”. In: *European Conference on Computer Vision (ECCV)*. Springer, pp. 69–84.